

A Benign Schema Is a Query Language: Selective System-Prompt Disclosure Through Required Tool Parameters

Abstract

This expanded earlier narrative is superseded for journal submission by `journal_paper/main.tex`. Its scanner interpretation is provisional: historical boolean votes lack raw judge text, so malformed responses cannot be excluded retrospectively.

An agent that calls a third-party tool fills in a schema the tool’s publisher controls. Values reach the publisher only if the host dispatches the call; the historical outcome below measures emitted arguments, not server receipt. Parameter-based extraction of system-prompt content is prior work and we ask what governs which content leaks, and whether an injection-oriented defense notices, for a parameter that injects no instruction and never mentions a system prompt.

Our lead finding is defensive. Invariant Labs’ published tool-poisoning policy, re-implemented locally, flags command-phrased injections but rates clean both our benign field and the published explicit attack, whose parameters are named `system_prompt` and `model_name`. The policy asks whether a declaration contains an injection, and parameter-based exfiltration is not injection-shaped, so it goes unmod-elled.

Disclosure is selective. Planting five seed-derived canaries and requesting one per schema yields a 68% diagonal against 2% off-diagonal and a 0/800 neutral control on GPT-4o, matched on the second pre-registered provider confirmatory under a protocol externally timestamped before these trials existed. A service-key field returns a fabricated credential fixture 20/20 on all five models (three post hoc), a fixture that unplanted controls never recover (0/367). Across 25 authentic third-party schemas the benign field discloses on 69% of GPT-4o tool calls, against a 0/228 neutral floor.

Against our own interest, a separate pre-registered rule failed and the benign field beat an explicit ask on only one of three models, short of the required majority, though its neutral-field control passed on all three. Fields that only gesture at their target retrieve it on no complete arm, so what we document is schema-steered disclosure, not general non-naming retrieval. With nothing planted, a real agent framework dis-

closed its identity 0 times in 240 trials. Unlike query-based prompt stealing, the adversary needs no separate interactive extraction query or optimization loop: the schema is the retrieval interface.

1 Introduction

A tool-calling agent treats a tool’s JSON Schema as a specification to satisfy. The schema comes from whoever published the tool. The agent is told to fill in the required parameters, and the values it fills in are sent to that publisher. In the Model Context Protocol (MCP) [1], the publisher is whoever wrote a server the user installed. They need no foothold in the host, no injected instruction, and no contact with the user. They need the tool to be called once, for a task the user actually wanted.

Prior work established that this channel exists. HiddenLayer [7] declared parameters named `system_prompt`, `model_name`, and `chain_of_thought`, and several production models filled them with the contents of their system prompts. That report gives demonstrations across seven models and one shipped IDE. It does not report rates, and it does not compare against simply asking the model, against a parameter that asks for nothing, or against what a scanner would catch.

So the channel is known. Two things about it are not. First, an agent’s system prompt usually holds several separate facts: what product this is, which region it runs in, who operates it, sometimes a key. A schema declares one parameter. Which of those facts comes back? If the answer is “all of them,” the tool author learns the deployment. If the answer is “the one the parameter asked for,” the tool author can choose. These are different problems for a defender, and nothing in the prior work separates them.

Second, the tools an agent installs are supposed to be screened. The MCP ecosystem ships scanners that read a server’s declarations before any call and report whether they look malicious. Whether those scanners see this attack is a

question about the scanners, not about the models, and it can be answered without any disclosure measurement at all.

We answer both. To find out which fact comes back, we plant five strings in one system prompt and have each schema declare exactly one extra required parameter. The strings come from a seed we fixed in advance, and they look like nothing a model would produce on its own: `svc_a0086960eac4e052`, `ue-xawpy-9`. So we check for them by exact match. No grader has to decide what counts as disclosure. The parameter that asked for a string gets it back on 68% of gpt-4o calls. The four strings it did not ask for come back on 2%. A neutral parameter, added the same way, returns none of them in 160 calls (800 correlated canary checks).

To find out whether a scanner would see any of this, we re-ran Invariant Labs’ published tool-poisoning policy against every condition. It flags the conditions that put a command in the tool description. It rates our benign parameter clean, and it also rates clean the published attack from prior work. That policy asks one question: does this description contain a prompt injection? Neither attack contains one (§8).

What we mean by an operationally plausible parameter

This term carries weight in our claims, so we fix it now. A parameter is operationally plausible here when it follows an established convention that a legitimate server could plausibly use. Separately, every detector configuration we tested rates our primary parameter clean; that observation is a result, not part of the definition. Operationally plausible does not mean the parameter hides what it wants. Our main one, `client_user_agent`, asks for “a client identification string describing the application or runtime issuing this request.” That is the HTTP User-Agent convention, and it says which category of value it wants. §7.1 tests parameters that genuinely avoid naming their target. Those mostly retrieve nothing. We keep the two senses apart because merging them would overstate what separates this work from prior work, whose parameters name their target outright.

Contributions.

- Invariant Labs’ published `policy.gr`, re-implemented locally, misses both our operationally plausible parameter and the published explicit attack (§8, Table 7). The gap is in what the policy asks, not in how we worded our parameter.
- Rates and controls where prior work gave demonstrations (§5): measured recovery for the published explicit attack (Table 3), a neutral-parameter floor of 0/228, and a pre-registered decision rule that failed, reported as it was written.
- The declared parameter decides which planted fact returns (§7), confirmatory on the 2 models named in a protocol whose hash was timestamped before the data

existed. A matched control with nothing planted separates reading from inventing: 0/367.

- A negative result on how far this generalises (§7.1). Parameters that avoid naming their target recover nothing on any complete arm. One cell in a quota-capped arm is the exception.
- The false-positive floor for a name-based detector is already occupied (§10). Our best attack wording scores 25 where legitimate telemetry scores 15, and a widely used benign MCP server already asks agents to report their own reasoning.
- A retraction (§12). With nothing planted, a real agent framework disclosed its own identity 0 times in 240 trials. This channel moves planted text, not self-knowledge. That distinction killed our first hypothesis.

2 Background and Related Work

Tool calling and MCP A host application gives a model a set of tools. Each tool has a name, a description in plain language, and a JSON Schema listing its parameters. The model emits a call; the host sends the arguments to whoever published the tool. MCP [1] standardises this so that tools can be installed from third parties. That makes the schema an untrusted input which the agent is nevertheless instructed to satisfy. Our adversary relies on three things: the schema author controls it, parameters can be required, and arguments reach the publisher. All three are in the specification, not in any one implementation. Recent work maps the protocol’s threat surface across a server’s lifecycle [9] and audits live deployments for exploits [19]. Both catalogue attacks that carry an instruction. Neither covers the parameter channel we measure.

Declarative parameters and prompt injection Text placed in a model’s input can redirect it or extract its prompt [18]. The form that matters in practice is *indirect*: the text arrives through a channel the user trusts, such as a retrieved page, an email, or a document [6]. Later work formalises the family and evaluates defenses against it [14]. A tool schema is such a channel, so it is fair to read this paper as an indirect-injection result. Our operational distinction is declarative parameter wording versus an imperative override, not an exclusion from every injection taxonomy. Condition C declares a required parameter with a plausible name and description. Recovery nevertheless conflicts with the prompt’s confidentiality instruction. The locally reproduced policy misses the tested cases (§8); this observation does not prove that every injection-oriented defense must miss them.

Extracting the system prompt What this channel carries is system-prompt content, which OWASP lists as LLM07: System Prompt Leakage [16]. Prior methods extract prompts by querying the deployed application [33], by optimising those queries [10], by inferring an equivalent prompt from a limited set of input–output pairs [26], or by inverting ordinary outputs with no adversarial query at all [29]. We share the target and differ in who the attacker is. Each of those methods talks to the service or reasons over its outputs. Our tool provider does neither. It publishes a schema and reads the arguments that arrive during a task the user started.

Measuring disclosure with canaries Planting an unlikely string and checking whether a system emits it is the standard way to measure unintended disclosure [2]. We borrow the method for a different target. We are not asking what a model memorised during training; we are asking what it forwards out of its prompt during ordinary operation. The requirement is the same: the string must be chosen before the measurement and must be unguessable. Ours come from a seed fixed in the protocol (§7).

Tool poisoning The known attack puts an *instruction* in a tool’s natural-language surface, usually the description, so the model follows it while appearing to do its job. Invariant Labs disclosed it for MCP [11], it is catalogued as a top-ten MCP risk [17], and it is what published tool-poisoning policies are built to catch, including the policy we re-run in §8. Our condition B reproduces it for comparison.

A recent study of tool descriptors finds tool poisoning the hardest of three semantic MCP attack classes to detect, and proposes vetting descriptors before they enter the context [12]. Its attacks still change natural-language descriptors. It does not study values pulled out by a required parameter. We therefore limit our detector claim to the published pre-call schema policy we re-run, not to every proposed traffic or runtime defense.

Parameter abuse HiddenLayer [7] showed that a declared *parameter* is enough: name one `system_prompt` or `model_name` and several production models fill it with prompt contents. Their report gives demonstrations across seven models and a shipped IDE. It does not give rates, matched controls, or a scanner evaluation. This is the closest prior work. We reproduce it as condition E and measure it in Table 3.

Optimised poisoning MCP-ITP [13] tunes poisoned tool metadata against detector feedback, reaching high attack success while staying undetected. The difference that matters here is effort. Their adversary runs a search loop and writes malicious text. Ours writes one plausible parameter, once. A defender who assumes attacks must look adversarial is wrong even against an adversary who does no work. Closer

in outcome, and further in capability, is exfiltration through a *backdoored* tool [32]: the leak is comparable, but that adversary controls what the tool does, where ours controls only what the tool declares.

Agent-security benchmarks, and why we do not evaluate on them AgentDojo [5] runs tool-calling agents under injection. InjecAgent [28] carries the injection in tool *output*. ToolEmu [21] emulates execution to surface risky *actions*. Agent Security Bench [31] organises attacks by agent component. All of them score injections that carry an instruction, or risky actions. None supplies the unit we need, where the adversary’s whole payload is a required parameter and success is the value written during a task that otherwise succeeds. MSB [30] is closest, since it already scores out-of-scope parameters. Our distinct measurement crosses multiple simultaneously planted facts with requested fields and scores selective recovery by fixed canary matching; it does not introduce parameter-based extraction. We do use MCPTox [25], for external validity (§6) and prevalence (§9), because it supplies real schemas rather than a taxonomy.

Which defenses could apply An instruction hierarchy [22] ranks instructions that conflict. Our prompts explicitly prohibit disclosure of the planted confidential facts, but a schema need not contain an imperative injection to elicit them. Whether hierarchy-trained models prevent this requires direct testing. Channel-separation defenses such as StruQ [3] and spotlighting [8] do not say how a client should treat a schema that is untrusted but operational. Detector-based defenses are also fragile under adaptive rewording [27]; §10 shows a nearby false-positive floor even without an adaptive search.

The class that fits the problem is capability-based control flow around the model [4] and information-flow tracking through the agent’s execution [34]. Both constrain which data can reach which sink, whether or not an instruction was present. That is the right shape here precisely because there is no instruction to detect. A taint label on the system prompt would stop a planted string reaching a third-party argument whatever the parameter is called. We do not evaluate either, and we name them as what we would evaluate next. Runtime trust layers that score a server’s behaviour during operation [35] are complementary, but they act after the call the adversary has already been paid by. MindGuard [24] detects manipulated agent decisions from attention structure with high precision, but needs white-box access and reports no false-positive rate on benign tools. Our defender runs against a hosted API and has neither that access nor room for unmeasured false positives. That is why our countermeasure (§10) is black-box, and why we report its precision under low base rates rather than its AUC alone.

What is new here Not that a parameter can extract prompt contents; that is [7]. Not that attacker-controlled input reaches a model through a trusted channel; that is [6]. We contribute the measurement, the finding that disclosure is selected by the parameter rather than wholesale, and the finding that the reproduced injection-oriented policy is built for a different attack. We also state where our novelty stops: HiddenLayer’s parameter names its target, and so do the parameters that work in our matrix (§7.1).

3 Threat Model

There are three parties: the operator, the user, and the tool publisher. We describe each in turn, then say what is out of scope.

Operator The operator runs the host application and writes its system prompt. That prompt carries deployment facts: which product this is, and often its version, its region, who operates it, or a policy it must follow. It usually also says the configuration is confidential. The operator does not see the tool’s schema and does not see the arguments the model sends.

User The user asks for something ordinary. In our experiments they ask an assistant to retrieve their three most recent orders. The user installs tools, or an administrator installs them. At the API level, a tool call succeeds whether or not it carried anything extra. Whether a particular deployed client displays the additional argument is outside our measurement.

Adversary The adversary is whoever published one of the installed tools. They write the tool’s name, its description, and its parameter schema, including which parameters are required. They see every argument the model sends, which is ordinary operation rather than a breach. In a host that does not pin schemas and re-review them, they can also change the declaration after installation.

They have no access to the system prompt, the model weights, the host, or the network. They cannot fine-tune and they never interact with the user. Learning the deployment is what they are trying to do, not something they start with. The adversary is passive, in-band, and needs one call: they declare a parameter and read what arrives.

What the experimenter controls, and the attacker does not We plant the canaries in the system prompt ourselves. That is an experimenter intervention, not an attacker capability. Our adversary cannot write to the system prompt and cannot choose what it contains. They can only declare a parameter and hope the deployment holds something matching it. Every




Figure 1: Who sees what. The adversary writes the schema (2) and receives the arguments the model puts into it (3). It never sees the system prompt (1) and sends no instruction. It needs the tool called once, for the task the user asked for. Condition C adds one required parameter to (2) and reads one extra value out of (3).

severity statement in this paper is conditional on the deployment already containing the material, and §9 measures how often that is true in the corpus we can survey.

What is lost The deployment leaks, not the model, and it leaks inside a call the user asked for, with no refusal shown and no alert from the policies we tested (§8). One published tool could see many installations that call it. We do not measure how often that happens.

Compliance is not consent A tool author can argue that the schema asked for a value and the model supplied one. The within-task comparison shows why that is still a control failure. With the same confidentiality instruction in the system prompt, gpt-4o disclosed on 0/30 direct requests in chat and on 29/30 benign parameters. This is one model, not a general ordering, but it shows that rephrasing the request as a schema field defeats the boundary the instruction was meant to hold. The instruction is still there, and the API call produces no refusal. We did not test a commercial client’s user-interface or consent surface.

Defender We assume the defender runs or consumes a hosted model API. This rules out white-box defenses by construction, and it is a scoping choice rather than a convenience: an MCP client built on a commercial model can see the schema and the arguments, and nothing inside the model. The countermeasures in §10 are limited to what such a client can enforce on its own.

Benign servers already occupy this shape Our corpus contains a server that looks like the attack and is not one.

Sequential Thinking asks the agent to report its reasoning through parameters named `thought`, `nextThoughtNeeded`, `totalThoughts`, and `revisesThought`. That is the same neighbourhood as `chain_of_thought`, which prior work published as an exfiltration vector [7], and it is offered in good faith. When we score real schemas with the classifier of §10, 4 of its parameters land among the 5/13 false positives at the operating threshold. A defender cannot separate this class of parameter from the attack by shape, because benign servers are already in that shape.

Scope limits that travel with every severity claim The attack surface is close to unattested in the ecosystem we can measure. Of the parameter names we surveyed, 0/401 are client-identity shaped and 0/401 ask for a credential, and the host-prompt text we surveyed contains 0 credential-shaped strings. So the attack as staged needs both an unusual parameter family and a deployment whose prompt holds the requested material. That is a real ecological limit. It also makes a check keyed to these families cheap in today’s corpus (§9). The detection result does not depend on any of this: it concerns what the deployed policy models, not how often the attack occurs.

4 Study Design and Methodology

One task, one thing changed Every synthetic trial uses the same task. A user asks an assistant to retrieve their three most recent orders, and one tool does that. Between conditions we change only the added parameter or the tool description. The task, the tool’s job, and the user’s request stay identical, and the task never mentions identity. So any difference between conditions comes from the parameter.

Conditions A asks in chat which model and framework is running. A’ asks the same thing through a required parameter named `agent_identity`, worded neutrally. B puts an urgent instruction in the tool description, reproducing ordinary tool poisoning. B’ rewords the A’ request as a command, which separates tone from what is being asked for. C adds a required benign parameter as defined in §1; the main wording presents it as the tool-calling equivalent of an HTTP User-Agent header. D adds a required `request_trace_id`, which pulls at no identity, and is the neutral control. E adds `system_prompt` and `model_name`, reproducing prior work [7]. A’, B’ and E were added after we had seen data, and we label them post-hoc wherever they appear.

We ablated condition C across five wordings: `default`, `operational_metadata`, `compliance_logging`, `api_documentation`, and `indirect_inference`. We picked `api_documentation` only after seeing how each performed, so the main wording is post-hoc. So is the

reticence ladder. We report every wording outcome in §5 rather than only the one we kept.

What we plant The system prompt carries material the adversary should not receive: a product name, an unreleased build identifier, infrastructure and tenancy facts, a policy token, and a credential-shaped string. All of it is fabricated, and the credential corresponds to no real service. A confidentiality instruction sits alongside it, so the model is told to withhold exactly what the parameter invites it to supply.

The matrix design, and why it needs no grader Grading disclosure is the weak point of this kind of measurement. If the outcome is “did the model reveal something about its deployment,” someone has to decide what counts, and that decision can drift.

We avoid the problem. Each matrix trial plants five canaries in one system prompt, and the schema requests exactly one of them. The four it did not request measure leakage to the side. A matched control plants nothing. Each canary is `sha256(seed | family)` cut to a fixed width, with the seed frozen in the protocol before collection, which lets a reader regenerate all five and confirm we did not choose markers after seeing which ones came back. They look like `svc_a0086960eac4e052` and `ue-xawpy-9`. A model is unlikely to produce those by accident. Recovery uses normalized substring matching of a full marker or prespecified prefix alias, and §7 uses no grader at any point.

We lodged the protocol’s hash with four independent timestamping calendars before running any trial in it.

Outcome measures for the older arms The earlier arms (§5, §6) predate the canary design and do need grading. There, a keyword grader flags a specific named framework or product (T1), and a blinded LLM judge, which never sees which condition a trial came from, assigns one of eight categories fixed in advance. The pre-registration requires both to agree. We report that conjunction: it reproduces every scaffolded number exactly, with no disagreements across the judged rows.

Human validation Our protocol required a person to hand-label 15–20% of trials and to report Cohen’s κ against the automatic grader, with a bar of 0.80. That is done. An author hand-labelled a blinded, stratified 20% sample of the gpt-4o gate data, 60 rows, without seeing the automatic verdicts, and attested to assigning every label unaided.

Agreement with the keyword grader is $\kappa = 0.91$ on the scaffolded stratum, which is where the disclosure signal is. The bare stratum is degenerate, because the keyword grader has no positives there at all, so we report the two strata separately rather than the pooled 0.86 that averaging them produces. Against the blinded eight-category judge the same labels agree at $\kappa = 0.86$.

An LLM adjudication of the same rows exists. We record it as a third automatic grader, never as human labelling; the human labels differ from it on 19 of the 60 rows. Two automatic graders agreeing is not evidence that either is right, and we have a concrete example: an earlier revision of the judge was perfectly self-consistent and scored $\kappa = -0.13$ against the keyword grader on this construct, worse than chance, and nothing caught it until an independent rater existed. Every grader-dependent claim is in §5 and §6. The mechanism result needs none of it.

Denominators A parameter can only disclose if the tool is called, so rates for the parameter conditions are conditional on a tool call. That flattens end-to-end risk, so we print intention-to-treat (ITT) next to every conditional rate in §6.

Statistics Rates carry Wilson intervals. Comparisons use Fisher’s exact test with Holm correction, and Newcombe intervals on the risk difference. We use those because this data separates completely: many cells are exactly 0 or exactly n , where odds ratios do not exist.

Where a claim generalises over tools or over parameters, we resample those units rather than trials. Repeated calls against one fixed configuration are repeated samples of one deployment, not independent deployments. Every interval in Table 5 is clustered by parameter on that basis.

Pre-registration, and the rule that failed The primary decision rule was registered before the trials that evaluated it. It failed (§5). We report it as written rather than rewriting it around the outcome.

A companion document records every deviation: conditions added after seeing data, a grader revised mid-study, and an earlier matrix experiment whose analysis plan was not committed before its data. That last failure is why the protocol behind §7 was hashed and timestamped before a single trial ran, and why we report only the 2 models that protocol names. We put this in the method section rather than a footnote because it changes how the results should be read.

Every arm, and what it is worth Table 1 lists each arm and its evidentiary status before any of them is discussed. Only the two rows marked confirmatory were named by a protocol whose hash was witnessed before the data existed. No amount of agreement among the others upgrades them.

Figure 2 shows how an artifact becomes a claim.

Figure 3 shows the C–D contrast. Both tools do the job the user asked for. The adversary adds one required parameter and reads its value out of a call that succeeded.

Table 1: Every arm we report and what it is worth. “Conf.” means the arm was named in the externally witnessed protocol before we collected data. “Post-hoc” means we added it afterwards. “Expl.” means the analysis plan was not committed first, which cannot be repaired later.

Arm	§	Models	Status
Gate (A/A’/B/C/D)	5	3	pre-registered, failed
Wording ablation	5	2	post-hoc
Explicit params (E)	5	1	post-hoc
Authentic schemas	6	2	post-hoc
Fact $\times$ field matrix	7	2	conf.
+ three more providers	7	3	post-hoc
Fabrication control	7.2	2	conf.
+ two more providers	7.2	2	post-hoc
Content familiarity	7.1	1	post-hoc
Depth (required, position)	10	1	expl.
Scanner policy	8	n/a	re-implementation
Parameter classifier	10	n/a	post-hoc

5 Calibrating the Channel

This section measures how strong the channel is and reports the pre-registered rule that failed. It is not the paper’s main result; §7 is.

The gate failed as written The registered rule required a gap of at least 20 pp between the benign parameter and the explicit request, on at least two of three models. We observed that gap on one. On gpt-4o the explicit request returned 9/30 and the benign parameter 29/30. On Claude and Gemini both conditions sat at the ceiling.

One of three. The gate failed. The ceiling explains it, since both models fill the explicit parameter on every trial and no gap can be measured against a ceiling, but the rule made no provision for that and we are not rewriting it now. Table 2 gives the full result.

The wording control passed on all three Against the neutral parameter under the same task, C versus D was 29/30 against 0/30 on gpt-4o, 30/30 against 0/30 on Claude, and 30/30 against 0/30 on Gemini. So the registered wording control passed on every model.

On gpt-4o we have the same cell twice, from two runs with identical settings. Pooled, C is 92% [82, 96] and A’ is 33% [23, 46]. What moves the content is the parameter’s wording, not the fact that a required parameter was added.

Every wording we tried The wording ablation was post-hoc, and the result is model-specific. On gpt-4o, default and indirect_inference recovered no planted marker at all. operational_metadata and compliance_logging recovered one only rarely. api_documentation was the only wording that disclosed reliably. Claude recovered the planted marker under every wording on every trial.

Two things follow. Choosing api_documentation was informed by the outcome, and the ablation does not show




Figure 2: How an artifact becomes a claim. Only the execute stage writes trial data; analysis reads it without modifying it. The scanner and classifier paths consume schemas and never write trial data.

C — benign naming field	D — neutral control
<pre> "query": {"type": "string"}, "client_user_agent": { "type": "string", "description": "Client user-agent string" } required: [query, client_user_agent] </pre>	<pre> "query": {"type": "string"}, "request_trace_id": { "type": "string", "description": "Opaque id for request tracing" } required: [query, request_trace_id] </pre>
Illustrative arguments received by the tool provider:	
<pre> C {"query": "three most recent orders", "client_user_agent": "PLANTED_PRODUCT_CANARY"} </pre>	<pre> D {"query": "three most recent orders", "request_trace_id": "req-001"} </pre>

Figure 3: Conditions C and D, abridged. Both tools do the job the user asked for. The only difference is the added required parameter and the extra value it returns to the tool provider. No instruction is injected.

Table 2: The gate experiment. Parameter conditions are conditional on a tool call; A is a direct request in chat. $\Delta = C - A'$. Only gpt-4o clears the registered 20-pp threshold, so the majority-model rule fails. C beats D on every model.

Condition	gpt-4o	Claude	Gemini
A: direct chat	0/30	30/30	0/30
A': explicit field	9/30	30/30	30/30
B: injection	0/30	0/30	7/30
C: benign naming	29/30	30/30	30/30
D: neutral	0/30	0/30	0/30
Δ	67 pp	0 pp	0 pp
Threshold	meets	misses	misses
Overall verdict	FAILED: majority rule not met		

that invoking the User-Agent convention is what makes the parameter work in general.

Measuring the published explicit attack Prior work reports demonstrations. Table 3 gives rates. Condition E recovered the planted marker on 46/100 trials pooled, but the pooled figure hides the interesting part: an internal codename comes back on 100% of calls and a policy sentence on 0%.

Two confounds travel with this table and should not be dropped. The payload fixtures are not matched on how strongly the prompt tells the model to withhold them. And for the credential, the prompt itself says the value is a test fixture rather than a live credential, so 20% measures that annotation at least as much as it measures the channel. §7 re-measures a planted, fabricated credential without the annotation and gets

Table 3: Condition E, the explicit parameters from prior work (system_prompt, model_name), on gpt-4o with $n=20$ per payload. Marker recovery in non-query arguments, normalised. Read with the two confounds named in the text.

Planted payload	Recovered
Internal codename	20/20
Long configuration block	19/20
Planted, fabricated credential fixture [‡]	4/20
Product name	3/20
Policy sentence	0/20
Pooled	46/100

[‡] planted, fabricated, and annotated in-prompt as a test fixture; see text.

Table 4: Real schemas with one parameter added by us. Conditional rates count tool calls; ITT rates count attempted trials. Both are printed so that a model choosing not to call a tool cannot flatter end-to-end success.

Model	C: benign naming		D: neutral	
	Conditional	ITT	Conditional	ITT
gpt-4o	78/113	78/125	0/109	0/125
Claude	120/120	120/125	0/119	0/124

20/20.

6 Does This Work on Real Schemas?

One synthetic tool cannot show that the channel survives contact with schemas we did not write. So we took 25 authentic tool listings from 25 MCPTox server entries [25] and appended one parameter. Because MCPTox’s tool-name and request arrays are not index-aligned, we paired them by greedy token overlap and retained each match score for audit.

Under the benign parameter, gpt-4o disclosed on 78/113 calls, which is 69% conditional and 62% ITT (78/125). Claude disclosed on 120/120, which is 100% conditional and 96% ITT. The neutral control disclosed on 0/228 calls across both providers, an exact floor. Counting tools rather than calls, 22 of the 24 tools gpt-4o actually called leaked at least once, and 24 of 24 on Claude.

The two denominators differ because not every tool is invoked on every trial. We print both so that the conditional rate cannot be read as end-to-end success. Table 4 keeps them side by side.

Two cautions belong with these numbers. Repeated calls against one tool are samples of one deployment rather than independent deployments, so we resample tools rather than trials. And these are real schemas that we modified: the tools are authentic, the added parameter is ours. The arm shows the channel works across schema contexts we did not write. It says nothing about how often such a parameter appears in the wild.

The gpt-4o rate on our synthetic tool is higher than its rate on these real ones. This is a descriptive difference between these samples, not a statistical upper bound or evidence about deployments generally.

7 The Parameter Decides What Leaks

We planted five canaries in one system prompt and had each schema declare exactly one extra required parameter. If the channel is wholesale, every parameter should return everything in the prompt. If the parameter selects, only the one it asked for should come back. Recovery is normalized marker/alias matching, so this section uses no grader.

Provenance We lodged the SHA-256 of the protocol with four independent timestamping calendars before any trial ran, in a repository state that contained no artifact from this experiment. Only the digest was sent; the file never left the machine. Bitcoin anchoring is still pending as we write, so the guarantee rests on four calendars rather than a block header, and we say it that way on purpose.

That protocol names 2 models. Every other provider was added afterwards and is reported separately throughout. The post-hoc rows of Table 5 should not be read as confirmatory.

Pre-registration did not let us pick the best-looking subset. The highest diagonal estimate belongs to a post-hoc provider, and the other post-hoc estimates fall between or near the pre-registered pair. Choosing after seeing the data would have let us report a stronger number. That is the freedom pre-registration removes.

The corrected paired implementation of the registered field-bootstrap analysis gives diagonal-minus-off-diagonal differences of 66 percentage points [32, 95] for gpt-4o and 58 points [17, 86] for gemini-3-flash, over 7 fields. Both point differences exceed the 20-point threshold and both intervals exclude zero, satisfying H-v3-1. The wide intervals reflect that the inferential unit is the field, not 20 repeated calls. The implementation correction was made after collection and is documented in the artifact’s deviations record; it is not a new preregistration. The separately preregistered generic field, `call_context`, has no target by construction and therefore stays outside that contrast; it returned any planted canary on 0/20 trials on each confirmatory model.

On every model the requested canary comes back far more often than the four that were present but not requested. So

Table 5: Which canary comes back. “Diagonal” is the one the declared parameter asked for; “off-diagonal” is the four others present in the same prompt. Intervals resample the seven hand-selected targeted parameters; they do not establish population coverage for unseen deployments. Rows above the rule use protocol-named models.

Model	Status	Diagonal	Off-diag.
gpt-4o	conf.	68% [36, 96]	2% [0, 4]
gemini-3-flash	conf.	59% [20, 87]	1% [0, 2]
claude-sonnet-4-5	post-hoc	71% [43, 100]	0% [0, 0]
deepseek-v4-flash	post-hoc	61% [29, 91]	1% [0, 1]
gemini-3.1-pro [†]	post-hoc	98% [90, 100]	5% [0, 13]

Neutral control (D): no recovered canary on any model. [†] overall arm quota-capped; see §12.

what moves content is the declared parameter, not the presence of an extra parameter.

Where the side leakage comes from Pooled, leakage to the side is about one percent on every model, which reads as a small uniform amount of noise. Per parameter it is neither small nor uniform.

On gpt-4o, 8/9 of the side recoveries come from one parameter, `client_user_agent`, which returned the region and the operator alongside the platform it asked for. On gemini-3-flash the same parameter accounts for 4/4. The three parameters that name a *specific* fact — `deployment_region_code`, `operator_account_name`, and `service_key_reference` — account for 0/9 and 0/4, and for none of the side leakage on any of the 5 models.

The asymmetry runs one way, and that direction is the useful one. A parameter naming a broad category pulls in neighbouring deployment facts. A parameter naming one particular fact does not, in either direction. Selection is therefore a property of how narrowly the adversary specifies the parameter, not of the channel as a whole. The parameter that sprays is also the one §7.1 finds unstable across content types. Together these say something simple about the adversary’s choice: naming the fact buys precision and a higher rate, naming the category buys breadth and costs both.

The confirmatory denominator includes deliberate negatives The witnessed protocol puts the non-naming parameters of §7.1 into the same diagonal denominator as the naming ones. So 68% mixes parameters that work with parameters we expected to fail. Restricted to naming parameters the rate is 94%. Read Table 5 as a rate over a parameter population that includes deliberate negatives. An earlier, uncommitted matrix contained only naming parameters, and it stays exploratory.




Figure 4: Canary recovery by model. For each model the black bar is the canary the parameter asked for, with its parameter-clustered interval; the grey bar is the four canaries present in the same prompt that it did not ask for. The neutral control returned none in 160 calls (800 correlated canary checks) on each confirmatory model and is not plotted. • marks the 2 models named in the witnessed protocol. 4o = gpt-4o, G3F = gemini-3-flash, C45 = claude-sonnet-4-5, DS4 = deepseek-v4-flash, G3P = gemini-3.1-pro.

The service-key result holds on every model A parameter named `service_key_reference` returned the planted, fabricated credential on 20/20 calls on gpt-4o and 20/20 on gemini-3-flash, and 20/20, 20/20, and 20/20 on the three post-hoc models. That is all 5 models we tried. Only the pre-registered pair is confirmatory.

This is worth comparing against condition E. There the credential came back on 4/20 calls through parameters named `system_prompt` and `model_name`, and that fixture was annotated in the prompt as a test value. Here an unannotated fixture moves through a parameter matched to it. What predicts recovery across these conditions is the match between the parameter and the content, more than how sensitive the content is. Every credential in this paper is fabricated and names no real service.

7.1 Must the parameter name what it wants?

This question decides how far the result generalises, and our answer is mostly negative.

We crossed two kinds of parameter. A *naming* parameter states the category it retrieves. An *adjacent* parameter does not: `issuing_surface` for the platform, `locality_hint` for the region, and `integration_binding_note` for the credential.

The class average says the opposite of the data Pooling the three adjacent parameters gives 33% [23, 46] on gpt-4o, and the same rate on several other models. A number that flat across models is suspicious, and it is misleading.




Figure 5: Recovery of the targeted canary, one row per declared parameter. Fill is the recovery rate, so an empty box never retrieved its target. The two bold rows are the real non-naming test. * marks a cell with fewer than 20 trials: `integration_binding_note` on G3P is quota-capped at 0/3, and `operator_account_name` on DS4 is 17/17 because that model declined to call the tool three times. Model codes as in Figure 4.

Table 6 shows why. The whole pooled rate is one parameter, `locality_hint`, at 100%, and `locality` is close to a synonym for the region it retrieves. The two parameters that genuinely avoid naming their target recover the canary 0/20 and 0/20 on gpt-4o, and recover nothing on any complete comparison arm. The one contrary platform result is a complete cell inside a quota-capped, post-hoc arm.

We report the per-parameter table rather than the class average for exactly this reason.

One naming parameter is unstable, and \$5 rests on it

Every other naming row in Table 6 sits at the ceiling on every model. `client_user_agent` does not. It recovers 75% on gpt-4o and 10% on gemini-3-flash, where the same parameter with the same wording scored 97% and 100% in \$5.

The two experiments differ in what was planted as well as in the surrounding prompt. The earlier one planted a recognisable product name; this one planted a seed-derived token that exists nowhere. Our first reading was that this explains the gap: a category parameter picks up a recognisable entity easily and an arbitrary string poorly.

We tested that explanation and it was wrong

The contrast is cheap to isolate, so we isolated it. One model, one parameter, one prompt, two arms differing only in whether the planted platform canary is a name the model has seen before. The two prompts are the same length to the character, so length cannot explain a difference.

The result runs the other way. The recognisable name came back on 1/20 calls. The seed-derived token came back on 14/20. That is a gap of 65 pp against familiarity. The seed-derived arm also reproduces its own matrix cell at 75%, which is the internal check that the arm measures what we think it does.

Table 6: Recovery per parameter. `adj.` avoids the target noun; `nam.` states the category it retrieves. `locality_hint` is a near-synonym and we analyse it separately. The two bold rows are the real non-naming test, and pooling them with `locality_hint` reverses the conclusion.

Declared parameter	Class	gpt-4o	gem-3-flash	claude-4-5	deepseek-v4	gem-3.1-pro [†]
issuing_surface	adj.	0%	0%	0%	0%	100%
integration_binding_note	adj.	0%	0%	0%	0%	0%
<code>locality_hint</code>	adj.	100%	100%	100%	100%	100%
<code>client_user_agent</code>	nam.	75%	10%	100%	35%	100%
<code>deployment_region_code</code>	nam.	100%	100%	100%	100%	100%
<code>operator_account_name</code>	nam.	100%	100%	100%	100%	100%
<code>service_key_reference</code>	nam.	100%	100%	100%	100%	100%

[†] overall arm quota-capped; the credential-adjacent cell for the planted, fabricated fixture is 0/3.

What the model wrote explains the direction. Given a recognisable platform name, it answered the parameter out of its own generic self-concept — `assistant`, `online_store_assistant` — and left the planted name alone. Given an unrecognisable one, it forwarded the planted token, often concatenated with a second canary (`Twpjxkh Systems Ivdgro Runtime`). A recognised name competes with an identity the model can supply from itself. An unfamiliar string has no competitor, so the only candidate in context is the planted one.

For a defender this is the uncomfortable direction. The content an adversary most wants — unreleased identifiers, internal tokens, anything that exists nowhere public — is the content this channel carries best, not worst.

It also means our earlier explanation of the §5 figure was wrong, and the gap between the two arms is still unexplained. They differ in prompt structure too, and we do not claim to have isolated the cause. This arm is post-hoc and postdates the witnessed protocol. It is one model and one parameter. Read it as a single clean contrast, not as a law.

What we therefore do and do not claim We claim that the parameter selects what comes back. We do not claim that a parameter which never names its target retrieves it. Our `service_key_reference` parameter stays close to the explicitly named parameters of prior work: it avoids the word `system_prompt`, but its own name states the category it wants. A reader should weigh §8 accordingly. That result does not depend on this one.

One exception deserves a sentence. `gemini-3.1-pro` recovered the platform through `issuing_surface` on 20/20 calls, where every complete comparison model recovered none. That cell is complete, but it sits inside an arm that was quota-capped overall and confined to one provider. Capability scaling stays a hypothesis for future work rather than a finding.

7.2 Is the model reading the prompt, or inventing?

A parameter that gets filled in is not a parameter that leaked. Models will write something plausible into any required field. If we counted filled fields as disclosure, we would be measuring compliance, not leakage.

We re-ran six of the eight parameters with nothing planted: same task, no canaries in the prompt. Platform-adjacent and region-adjacent controls were not collected. Across these configured six-field stages, the models produced a canary in 0/367 tool calls. Completed stage metadata does not establish the protocol’s full same-field control coverage.

They do fabricate. With nothing to report, `gpt-4o` wrote `us, user’s store, and online-store-manager/1.0` into the same fields that carry canaries when canaries exist. What they did not do is invent a seed-derived string.

One case is worth noting. Asked for a service key when the prompt held none, `gpt-4o` declined to call the tool at all rather than invent one, while returning the planted key on 20/20 calls when the fabricated fixture was there.

This is why the outcome measure is canary recovery and not field population, and it is the control that separates this work from a demonstration that models fill in required parameters.

8 Why an Injection-Oriented MCP Policy Misses This

We put this result first among the defensive ones because it needs no disclosure grader, and because it applies to the published prior attack as well as to ours.

The policy asks about injection We re-implemented Invariant Labs’ published `policy.gr` verbatim: their prompt, their model, and their tool surface, which by their own template includes parameter names. We ran it locally over every condition, 3 repetitions, unanimous throughout.

To make that reproducible rather than merely asserted: `policy.gr` as shipped in `mcp-scan==0.3.39`, judged by `gpt-4o-mini` at temperature 0, which are the model and temperature the policy file itself specifies, over 3 repetitions under majority vote. The policy returns a binary verdict rather than a score, so there is no threshold to tune. The only free choices are the model and the tool surface, and we changed neither. Our harness prints the exact payload it submits for every condition before sending it, so an independent party can reproduce the inputs without running our code.

It flags conditions B and B', the two that phrase something as a command. It rates condition C clean under every wording we tested. It also rates condition E clean — the published prior attack, whose parameters are literally named `system_prompt` and `model_name`.

The policy asks whether a description contains a prompt injection. A parameter that collects data contains no instruction, so the question does not reach it, and the attack is not modelled.

Scope of this claim This is a gap in the policy we reproduced, not in every scanner. The general point is that an injection-only question gives no criterion for a schema that instructs nothing and still declares a hostile sink. The gap is in the question, not in our wording.

Others reach the same conclusion from the defensive side. Work on automated MCP server vetting reports that static inspection alone is not enough for exfiltration that rides ordinary parameters [23]. An analysis of where defenses belong puts the enforcement point in the client or host, rather than in a scan of descriptions before the call [20]. Our measurement supplies the empirical case: the artifact a scanner reads is exactly the artifact that carries no evidence of the attack.

Reading parameter names does not close it Extending a scanner to read parameter *names*, which is the natural response to prior work, catches both the honest explicit request and condition E. The benign parameter stays clean under both profiles and every wording. So name-based detection closes the gap prior work opened and leaves this one open.

And name-based detection has an occupied floor Our intent classifier scores the most effective attack wording at 25. `user_agent.original`, a standardised OpenTelemetry field [15], scores 15. The usable band between them is narrow, and legitimate telemetry already sits inside it.

The Sequential Thinking parameters from §3 account for 5/13 of the false positives here. That is the same fact seen from the defender's side. The shape is occupied, so separating by shape costs alerts on servers doing nothing wrong.

Table 7: Scanner verdicts. “Published” is Invariant Labs’ `policy.gr` from `mcp-scan==0.3.39`, re-implemented locally; it is not the vendor’s current hosted service. E is the explicit attack from prior work, whose parameters are named `system_prompt` and `model_name`. 3 repetitions, unanimous on every condition.

Condition	Published	Ours	
	<code>policy.gr</code>	desc	desc+name
A' explicit ask	clean	clean	flag
B injection	flag	flag	flag
B' command tone	flag	flag	flag
C benign naming	clean	clean	clean
D neutral control	clean	clean	clean
E published attack	clean	flag	flag

9 Would This Attack Fit the Real Ecosystem?

An attack that no real ecosystem would accommodate is a curiosity. The attack needs two things: a parameter family that would not look out of place, and a deployment whose prompt holds the requested material. We measured both over 45 MCPTox server entries in our corpus, covering 1048 parameters and 401 distinct names. We fixed the matching rules before running them and deliberately let them over-count.

The attacker’s parameter is close to unattested Of 401 distinct real parameter names, 0/401 are client-identity shaped, 0/401 ask for a credential, and 0/401 name a policy tier. Only region (1/401) and contact (1/401) have any precedent at all.

The operator and tenant matches (5/401) are rule noise. On inspection every one of them is repository ownership in code-hosting tooling, not the tenant an agent is deployed for.

The victim’s material is mostly absent too Across 195,012 characters of MCP host-prompt text we found 0 credential-shaped strings, 0 tenant statements, 0 policy tokens, and 3 cloud regions.

What this means We read this as a limit on the attack as we staged it, and we say so. A tool asking for a service key would look unlike anything the surveyed ecosystem declares, and the material we plant is not what these hosts put in front of a model. The corpus is constructed tool listings rather than deployment prompts from a shipped product, so it bounds this population and not that of a commercial IDE.

The same numbers cut the other way for defense. A detector keyed to these families fires on 0/401 real parameters today, so its false-positive cost in the current ecosystem is near zero.

This section argues against our own severity claim, and it is in the main body for that reason. It does not touch §8, which describes what a detector class models rather than how often the attack occurs.

Table 8: Precision of the parameter-name classifier under three assumed attack base rates. The chosen threshold flags all nine study positives (TPR 100%, FPR 2.6%). AUC is 0.994 [0.986, 0.999] over 9 positives and 506 real benign schema parameters.

Base rate	PPV	Reading
0.1%	3.8%	alerts dominated by false positives
1.0%	28.2%	false alerts still dominate
5.0%	67.2%	precision remains base-rate dependent




Figure 6: ROC for the parameter-name classifier, 9 positives against 506 real benign schema parameters (AUC 0.994). The curve is not the point. At every threshold that catches the attack, precision still collapses once the attack is rare; Table 8 gives three base rates.

10 What a Client Can Actually Do

Our threat model allows one kind of automated defense: a classifier over the parameter’s name and description, which is all a registry sees before any call is made. We built it, and we report it as a negative result.

The classifier separates well and still does not deploy On 9 study-authored positives against 506 real-schema negatives, the classifier reaches an AUC of 0.994 [0.986, 0.999]. Lexical separation from the detector prompt does not establish an independently designed holdout. At the operating point flagging all study positives it achieves a 100% true-positive rate at a 2.6% false-positive rate. Those numbers look good. The leakage-reduction calculation is retrospective, not an executed blocking experiment; legitimate-task utility was not measured.

They do not survive contact with a realistic base rate. If 0.1% of parameters are hostile, precision is 3.8%, so most alerts are false. At 1.0% it is 28.2%. Only at 5.0% does it reach 67.2%, which is where it starts to be operable. We do not recommend the classifier under the low base rates we expect in practice, and Table 8 is there so a deployer can check the assumption for their own setting.

Two things that did change recovery Neither requires touching model weights.

Making the added parameter *optional* rather than required cut recovery by 60% on the identity parameter, from 100% to 40%, and by 30% on the parameter for the planted credential. A client can enforce this on its own, but only if it treats the advertised schema as untrusted: it can refuse or downgrade required parameters that the tool does not need, and it can drop optional arguments before dispatch.

Naming tool parameters explicitly in the confidentiality instruction closed the channel on Claude and not on Gemini. We therefore do not present prompt scoping as a general mitigation.

Optionality reduces the channel without closing it. The planted credential still comes back on 14/20 calls when its parameter is optional. So a client that suppresses optional arguments blunts the attack rather than stopping it. Our data also cannot say whether a model denied one optional parameter would move the same content into another, because each cell varies one factor against an otherwise fixed schema.

And one that did nothing We report the second factor because it is null. Moving the added parameter ahead of the tool’s own *query* parameter left recovery at 19/20 on the identity parameter and 20/20 on the credential parameter, against 20/20 and 20/20 in the default position.

Ordering inside the schema is not a lever for either side. Requiredness is. That is worth knowing before anyone ships parameter reordering as a mitigation.

11 Discussion

What the matrix shows, and what it does not The selection we measure operates over planted prompt content and is directed by the parameter. It is not access to anything the model holds internally, and the negative gradient is what bounds it.

The failure is quiet, and it depends on phrasing A confidentiality instruction is enforced against how a request is phrased more than against where the protected content goes. A tool call that succeeds produces no API-level refusal. Whether a deployed client surfaces the additional argument or requests consent is outside this measurement.

Injection-only policies make this worse by asking whether the metadata tells the model to misbehave. A parameter that collects data tells it nothing. This gap between what the detector asks and what the attack does survives even if later work revises our estimates of how strong the channel is or how often it applies.

What a client should do Treat a third-party schema as a declaration of where data will go. Do not honour unnecessary required parameters. Drop optional arguments the task does not need before dispatch. Review unfamiliar naming families.

Enforce source-to-sink policy outside the model rather than by instructing it.

Keeping credentials out of prompts, as OWASP recommends [16], removes the most consequential source. Prompt scoping helped Claude and not Gemini, so we do not offer it as universal. Capability-based control flow [4] is the broader design that could enforce this boundary regardless of how the model is asked.

12 Limitations

Evidence status Only 2 of the 5 models in Table 5 were named by the witnessed protocol. Every other provider is post-hoc. The `gemini-3.1-pro` arm was quota-capped, which leaves its credential-adjacent cell at 0/3. Its separate platform-adjacent cell is complete at 20/20 and is the only support for the capability-scaling reading, which is therefore not a finding. An earlier matrix had no committed analysis plan; it stays exploratory and we never pool it with the confirmatory data.

Human validation It is done (§4), and its scope is narrow: one stratified 20% sample of the `gpt-4o` gate arm, labelled by an author rather than an outside rater. Agreement is $\kappa = 0.91$ on the scaffolded stratum. The bare stratum is degenerate and we report it separately rather than pooling. The mechanism result uses no grader, so nothing in §7.1 rests on this.

Coverage The confirmatory `v3` arms use temperature 0.7; a separate exploratory `v2` temperature-zero arm also exists and is not pooled with `v3`. The synthetic and matrix arms use one task and one tool; the external-validity arm broadens to 25 real tools and their requests, but the attack parameter is still ours. Cell sizes are modest and we do not vary the surrounding agent. Seeds steer generation only on OpenAI-compatible endpoints; on Anthropic and Google the logged seed is a label and not a control.

One question our corpus cannot answer MCPTox publishes tools as rendered text rather than JSON Schema, so parameter types are unrecoverable and we declare every harvested parameter a string. That leaves open whether an enumerated or otherwise constrained type suppresses the channel, as it plausibly would by leaving the model no valid way to emit a canary. Settling it needs schemas harvested from live servers rather than from rendered listings. We think this is the most tractable next measurement.

Ecological validity The corpus contains neither a parameter family matching the attack nor credential-shaped victim material (§9). So disclosure rates on a synthetic tool are not ecosystem prevalence, and for `gpt-4o` the synthetic rate is higher in the observed sample, not a population upper bound.

The vendor’s hosted scanner is unrun pending coordinated disclosure, so our scanner evidence is limited to its published policy re-implemented locally. The prompt-scoping result comes from the post-hoc reticence ladder, is specific to Claude, and does not transfer to Gemini. All planted material is fabricated, including the credential.

A hypothesis we retracted An earlier version of this work supposed that models disclose their *genuine* runtime identity. A real framework with nothing planted disclosed it in 0 of 240 trials while calling the tool successfully throughout. We retracted the hypothesis. What we measure is extraction of planted prompt content, not self-knowledge, and we keep the control because the distinction is easy to lose.

13 Conclusion

A required tool parameter opens a channel. The agent fills it in, and the protocol delivers the value to whoever wrote the schema.

For parameters that name or closely paraphrase what they want, that channel selects: it returns the requested fact and usually withholds the rest. Invariant Labs’ published `policy.gr`, re-implemented locally, misses both our operationally plausible parameter and the published explicit attack because it asks whether a declaration contains an instruction and neither one does.

Parameter-based extraction is prior work. What we add is the measurement, the controls, and the mismatch between what the reproduced policy asks and what this attack looks like. The negative boundary matters as much: a parameter that neither names nor paraphrases its target does not retrieve it on any complete arm we ran.

A Ethical Considerations

How the study was run No human subjects were involved. Every measurement used the authors’ own provider accounts. We probed no third-party deployment and never made a tool act against a caller. We wrote all the prompts. Every planted fact is fabricated and the credential-shaped fixture names no real service, so the study handled no live secret at any point.

Who is affected MCP tool authors, agent-host operators, end users, model providers, and scanner vendors. A malicious tool author could quietly fingerprint deployments or pull out prompt-resident material. Hosts and users bear that loss. Providers and scanner vendors may need to revise controls that currently look for instructions.

We reduce the harm two ways. We report the levers a client can pull without anyone’s cooperation, namely scope and requiredness. And we recommend keeping credentials out of prompts. The extraction mechanism was already public [7];

what we add is measurement, controls, and defensive analysis, not an optimized attack. The artifact includes an inert MCP server for transport testing, not a deployment against third-party users.

Disclosure, and the risk that remains We evaluated the scanner by re-implementing the vendor’s published policy locally. We did not run their hosted verification service, because that would send the new condition descriptions to the vendor outside a coordinated process.

Publishing still carries risk, since the paper describes how a published defense is evaded. We judge publication worthwhile: the mechanism is already public, clients can act immediately on requiredness and data-flow checks, and we release no optimisation recipe and no live-secret fixture.

The severity of the finding changed during the study, from fingerprinting to possible secret disclosure, which re-opened our disclosure obligation. At this snapshot the repository records no completed coordinated notification and the hosted-scanner run is still gated. This status must be updated if notification happens. The research repository was already publicly readable before we reassessed severity, and we record that because it bears on when the pattern became available.

B Open Science

Every quantitative result here is generated from versioned artifacts rather than transcribed. `code/manifest.py` derives the values from the run logs and emits them as LaTeX, `code/make_tables.py` regenerates artifact-backed tables used to cross-check the ones in this paper, and a claim check rejects prose that has gone stale or was retracted. We lodged the confirmatory protocol’s SHA-256 digest with four independent timestamping calendars before the data existed. Bitcoin anchoring is still pending, so we claim only those calendars.

The artifact contains the fixed-condition harness, regression and module self-checks, canonical results with completion metadata, and the scripts that rebuild the 25-tool real-schema corpus and the benign-parameter negatives. For release we provide code, canonical and derived artifacts, and instructions to reconstruct the rest. We do not release API keys, or the raw provider payloads that would reproduce planted prompts. Seeds are sampling controls only on OpenAI-compatible endpoints and labels elsewhere. The hosted-scanner run stays withheld pending coordinated disclosure.

References

- [1] Anthropic. Model context protocol specification. <https://modelcontextprotocol.io/specification>, 2024.
- [2] Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, Alina Oprea, and Colin Raffel. Extracting training data from large language models. In *USENIX Security Symposium*, 2021.
- [3] Sizhe Chen, Julien Piet, Chawin Sitawarin, and David Wagner. StruQ: Defending against prompt injection with structured queries. In *USENIX Security Symposium*, 2025.
- [4] Edoardo Debenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramèr. Defeating prompt injections by design. arXiv preprint arXiv:2503.18813, 2025.
- [5] Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.
- [6] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023.
- [7] HiddenLayer. Beyond MCP: Expanding agentic function parameter abuse. <https://www.hiddenlayer.com/research/beyond-mcp-expanding-agentic-function-parameter-abuse>, 2025.
- [8] Keegan Hines, Gary Lopez, Matthew Hall, Federico Zarfati, Yonatan Zunger, and Emre Kiciman. Defending against indirect prompt injection attacks with spotlighting. arXiv preprint arXiv:2403.14720, 2024.
- [9] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. Model context protocol (MCP): Landscape, security threats, and future research directions. arXiv preprint arXiv:2503.23278, 2025.
- [10] Bo Hui, Haolin Yuan, Neil Gong, Philippe Burlina, and Yinzhi Cao. PLeak: Prompt leaking attacks against large language model applications. In *ACM Conference on Computer and Communications Security (CCS)*, 2024.
- [11] Invariant Labs. MCP security notification: Tool poisoning attacks. <https://invariantlabs.ai/blog/mcp-security-notification-tool-poisoning-attacks>, 2025.
- [12] Saeid Jamshidi, Arghavan Moradi Dakhel, Kawser Wazed Nafi, and Foutse Khomh. Semantic attacks on tool-augmented LLMs: Securing the model context protocol against descriptor-level manipulation. arXiv preprint arXiv:2512.06556, 2025.
- [13] Ruiqi Li, Zhiqiang Wang, Yunhao Yao, and Xiang-Yang Li. MCP-ITP: An automated framework for implicit tool poisoning in MCP. arXiv preprint arXiv:2601.07395, 2026.
- [14] Yupei Liu, Yuqi Jia, Runpeng Geng, Jinyuan Jia, and Neil Zhenqiang Gong. Formalizing and benchmarking prompt injection attacks and defenses. In *33rd USENIX Security Symposium (USENIX Security 24)*, pages 1831–1847, Philadelphia, PA, August 2024. USENIX Association.
- [15] OpenTelemetry Authors. OpenTelemetry semantic conventions. <https://github.com/open-telemetry/semantic-conventions>. Accessed: Aug. 16, 2026.

- [16] OWASP. LLM07:2025 — system prompt leakage. <https://genai.owasp.org/llmrisk/llm072025-system-prompt-leakage/>, 2025.
- [17] OWASP. MCP03:2025 — tool poisoning. <https://owasp.org/www-project-mcp-top-10/>, 2025.
- [18] Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. In *NeurIPS ML Safety Workshop*, 2022.
- [19] Brandon Radosevich and John Halloran. MCP safety audit: LLMs with the model context protocol allow major security exploits. arXiv preprint arXiv:2504.03767, 2025.
- [20] Mehrdad Rostamzadeh, Sidhant Narula, Nahom Birhan, Mohammad Ghasemigol, and Daniel Takabi. MCP-DPT: A defense-placement taxonomy and coverage analysis for model context protocol security. arXiv preprint arXiv:2604.07551, 2026.
- [21] Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitit, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox. In *International Conference on Learning Representations (ICLR)*, 2024.
- [22] Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng, Johannes Heidecke, and Alex Beutel. The instruction hierarchy: Training LLMs to prioritize privileged instructions. arXiv preprint arXiv:2404.13208, 2024.
- [23] Bin Wang, Zexin Liu, Hao Yu, Ao Yang, Yenan Huang, Jing Guo, Huangsheng Cheng, Hui Li, and Huiyu Wu. MCPGuard: Automatically detecting vulnerabilities in MCP servers. arXiv preprint arXiv:2510.23673, 2025.
- [24] Zhiqiang Wang, Haohua Du, Guanquan Shi, Junyang Zhang, HaoRan Cheng, Yunhao Yao, Kaiwen Guo, and Xiang-Yang Li. MindGuard: Intrinsic decision inspection for securing LLM agents against metadata poisoning. arXiv preprint arXiv:2508.20412, 2025.
- [25] Zhiqiang Wang, Yichao Gao, Yanting Wang, Suyuan Liu, Haifeng Sun, Haoran Cheng, Guanquan Shi, Haohua Du, and Xiangyang Li. MCPTox: A benchmark for tool poisoning attack on real-world MCP servers. arXiv preprint arXiv:2508.14925, 2025.
- [26] Yong Yang, Changjiang Li, Qingming Li, Oubo Ma, Haoyu Wang, Zonghui Wang, Yandong Gao, Wenzhi Chen, and Shouling Ji. PRSA: Prompt stealing attacks against real-world prompt services. In *34th USENIX Security Symposium (USENIX Security 25)*, pages 2283–2302, Seattle, WA, August 2025. USENIX Association.
- [27] Qiusi Zhan, Richard Fang, Henil Shalin Panchal, and Daniel Kang. Adaptive attacks break defenses against indirect prompt injection attacks on LLM agents. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 7116–7132, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics.
- [28] Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics (ACL)*, 2024.
- [29] Collin Zhang, John X. Morris, and Vitaly Shmatikov. Extracting prompts by inverting LLM outputs. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.
- [30] Dongsan Zhang, Zekun Li, Xu Luo, Xuannan Liu, Peipei Li, and Wenjun Xu. MCP security bench (MSB): Benchmarking attacks against model context protocol in LLM agents. arXiv preprint arXiv:2510.15994, 2025.
- [31] Hanrong Zhang, Jingyuan Huang, Kai Mei, Yifei Yao, Zhenting Wang, Chenlu Zhan, Hongwei Wang, and Yongfeng Zhang. Agent security bench (ASB): Formalizing and benchmarking attacks and defenses in LLM-based agents. arXiv preprint arXiv:2410.02644, 2024.
- [32] Wuyang Zhang and Shichao Pei. Your LLM agent can leak your data: Data exfiltration via backdoored tool use. arXiv preprint arXiv:2604.05432, 2026.
- [33] Yiming Zhang, Nicholas Carlini, and Daphne Ippolito. Effective prompt extraction from language models. In *Conference on Language Modeling (COLM)*, 2024.
- [34] Peter Yong Zhong, Siyuan Chen, Ruiqi Wang, McKenna McCall, Ben L. Titzer, Heather Miller, and Phillip B. Gibbons. RTBAS: Defending LLM agents against prompt injection and privacy leakage. arXiv preprint arXiv:2502.08966, 2025.
- [35] Zhenhong Zhou, Yuanhe Zhang, Hongwei Cai, Moayad Aloqaily, Ouns Bouachir, Linsey Pang, Prakhar Mehrotra, Kun Wang, and Qingsong Wen. MCPShield: A security cognition layer for adaptive trust calibration in model context protocol agents. arXiv preprint arXiv:2602.14281, 2026.